

DS 6050 Final Project Sources

Transfer learning and fine-tuning with pre-trained MobileNetV2 network

1. Configure dataset for performance

- [`tf.data.Dataset.prefetch`](#): decouple time when data produced from time when data consumed
 - Prefetch and preprocess next set of inputs as model training on previous set
 - Background thread and internal buffer to prefetch elements from input dataset ahead of request
 - Num of elts to prefetch \geq num of batches consumed by a single training step
 - Manually tune or use `tf.data.AUTOTUNE`: prompts `tf.data` runtime to tune value dynamically at runtime

2. Data augmentation

- [`tf.keras.layers.RandomFlip`](#): randomly flips images during training
 - Arguments
 - `mode`: string indicating which flip mode to use (e.g., "horizontal" for left-right)
 - `seed`
- [`tf.keras.layers.RandomRotation`](#): randomly rotates images during training
 - Arguments
 - `seed`
 - `factor`: float represented as fraction of 2π or tuple of size 2 for lower and upper bound for rotating clockwise and counter-clockwise (if one float then the same for both)
- [`tf.keras.layers.RandomContrast`](#): preprocessing layer that randomly adjusts contrast during training
 - Arguments
 - `factor`: positive float represented as a fraction of value or tuple of size 2 for lower and upper bounds
 - When a single float then randomly pick between $[1.0 - \text{lower}, 1.0 + \text{upper}]$
 - `seed`

3. Rescale pixel values

- [`tf.keras.applications.mobilenet_v2.preprocess_input`](#): returns preprocessed `numpy.array` or `tf.Tensor` with input pixel values scaled between -1 and 1

- Alternative: `tf.keras.layers.Rescaling(1./127.5, offset = 01)`

4. Create base `MobileNet` model from pre-trained convolution network

- `MobileNetV2` model trained on the ImageNet dataset
- Pick layer of `MobileNetV2` to use for feature extraction
 - Use last layer before flattening → **bottleneck layer** (more generality vs. final / top layer)
 - Specify `include_top = False` to remove last classification layer(s) on "top"
- `next(iterator[, default]):` return the next item from the iterator

5. Feature extraction

- **5a.** Freeze convolutional base model → set `layer.trainable = False` (or `base_model.trainable = False` for all layers) to prevent weights from being updated during training
- **5b.** Add layers to allow for classification of images
 - Average over 7x7 spatial locations with [tf.keras.layers.GlobalAveragePooling2D](#) (global average pooling for spatial data)
 - Arguments
 - `keepdims=False`: whether to keep spatial dimensions → default `False` ∴ output: 2D tensor with shape `(batch_size, channels) = (32,`
 - Add `tf.keras.layers.Dense` to convert features into a single prediction per image
 - Use activation function "softmax" to return probabilities image is of a particular class for 18 classes (i.e., places)
- **5c.** Build a model by chaining together data augmentation, rescaling, `base_model`, feature layers using the Keras Functional API
 - Note: `training = False` for `base_model` since model contains a `BatchNormalization` layer → runs in inference mode and does not update its mean and variance stats
 - [Keras functional API](#): create more flexible models than `tf.keras.Sequential` API
 - Builds a graph of layers since deep learning model (usually) as directed acyclic graph (DAG) of layers
 - Create input node: `inputs = keras.Input(shape = (784,))`
 - Contains info about shape, dtype of input data
 - Create new node in graph of layers by calling a layer on inputs object: `dense = layers.Dense(64, activation = "relu") \ x = dense(inputs)`

- Layer call draws an arrow from inputs to new layer → pass inputs to `dense` layer and get `x` as the output
- Add another hidden layer: `x = layers.Dense(64, activation = "relu")(x)`
 - Pass `x` to next hidden layer and store output as `x`
- Add output layer: `outputs = layers.Dense(10)(x)`
- Create a model: `model = keras.Model(inputs = inputs, outputs = outputs)`
- Training, evaluation, inference the same as `Sequential` models

6. Compile the model

- `model.compile(optimizer, loss, metrics)`
 - optimizer: [`tf.keras.optimizers.Adam\(learning_rate = base_learning_rate\)`](#) → implements Adam algorithm, a stochastic gradient descent method based on adaptive estimation of first- and second-order moments
 - Higher computational efficiency, lower memory requirements
 - loss: `"sparse_categorical_crossentropy"` → use for two or more classes when labels are integers
 - [`tf.keras.losses.SparseCategoricalCrossentropy`](#)
 - metrics: [`\['accuracy'\]`](#) → (num times preds = labels) / count

7. Train the model

8. Fine tuning

- Only training a few layers on top of `MobileNetV2` base model → can increase performance to train weights of top layers of pre-trained `base_model` to force more generic params to associate with specific dataset
 - Only fine-tune small num of top layers (not whole `MobileNetV2` model) → higher up (later) layers more specialized ∴ more specific to dataset on which model trained
- **8a** Un-freeze the top (i.e., later) layers of the `base_model` (the `MobileNetV2` instance) but keep bottom (i.e., earlier) layers untrainable
- **8b.** Re-compile the model → use a lower learning rate to prevent model from overfitting quickly
 - [`tf.keras.optimizers.RMSprop`](#): RMSprop algorithm maintains a moving (discounted_ average of square of gradients → divides gradient by root of average

- **8c.** Continue training the model → use argument `initial_epoch` to start training again based on training model to convergence previously
 - `initial_epoch = history.epoch[-1]` → begin at last epoch of first round of training (when all `base_model` params frozen)

Model Checkpoints:

- [`tf.keras.callbacks.ModelCheckpoint`](#): callback to save keras model or model weights at some frequency
 - Use to save the weights that yield the maximum validation accuracy since model may begin to overfit to the training data ∴ the last epoch may not have the weights that yield the best validation accuracy
 - Access the weights: `model.load_weights(checkpoint_path)`
- [`tf.keras.callbacks.EarlyStopping`](#): stop training the model when a monitored metric has stopped improving
 - `patience`: number of epochs without improvement to wait until the training process stops

Saving and Loadings Models

- Save just the weights will callbacks (i.e., `tf.keras.callbacks.ModelCheckpoint`) or [save the entire model](#) as a `SavedModel`
 - Make a directory: `!mkdir -p saved_model`
 - Save the model using the directory
`path: model.save('saved_model/my_model')`
 - Load a new instance of the model: `new_model = tf.keras.models.load_model('saved_model/my_model')`

Tensorflow on Rivanna

- TensorFlow on Rivanna: <https://www.rc.virginia.edu/userinfo/rivanna/software/tensorflow/>
- UVA software containers (creating singularity containers) and docker hub info: <https://www.rc.virginia.edu/userinfo/rivanna/software/containers/>
- Docker images on Rivanna: <https://www.rc.virginia.edu/userinfo/howtos/rivanna/docker-images-on-rivanna/>
- `uvarc/rivanna-docker`: <https://github.com/uvarc/rivanna-docker>
- Explanation for code `singularity pull docker://uvarc/tensorflow:2.8.0` to create local copy of container image in a local directory: <https://hub.docker.com/r/uvarc/tensorflow>

Additional sources for a different approach (when data is stored with a directory structure s.t. images for each class are in separate folders)

- From the DEEPLIZARD website and YouTube channel
 - [MobileNet Image Classification With TensorFlow's Keras API](#) → default size for `MobileNet`: (224, 224)
 - [Process Images For Fine-Tuned MobileNet With TensorFlow's Keras API](#)
 - [Fine-Tuning MobileNet On Custom Data Set With TensorFlow's Keras API](#)

Organizing and Correctly Labeling Files

- Search for matches with regex using `re.search`: <https://stackabuse.com/python-check-if-string-contains-substring/>
- Multiple if else conditions for conditional col: <https://stackoverflow.com/questions/48569166/multiple-if-else-conditions-in-pandas-dataframe-and-derive-multiple-columns>
- Pandas - dataframe groupby - how to get sum of multiple columns (using `agg` and a `lambda` function): <https://stackoverflow.com/questions/46431243/pandas-dataframe-groupby-how-to-get-sum-of-multiple-columns>
- Counting the number of files in a directory with Linux using `ls | wc -l`: <https://devconnected.com/how-to-count-files-in-directory-on-linux/>
- How to convert two columns of a dataframe to a dictionary with `dict(zip(df.col1, df.col2))`: <https://cmdlinetips.com/2021/04/convert-two-column-values-from-pandas-dataframe-to-a-dictionary/>
- Listing files in a directory with `os.listdir('dir_name')`: <https://www.codegrepper.com/code-examples/python/pandas+list+files+in+directory>
- Using the `os` module: <https://docs.python.org/3/library/os.html>
- Copying files in Python with the `shutil` module: <https://stackoverflow.com/questions/123198/how-to-copy-files>
- Using the `shutil` module: <https://docs.python.org/3/library/shutil.html>
- Dictionary comprehensions: <https://www.datacamp.com/community/tutorials/python-dictionary-comprehension>

- Create multiple subdirectories in Linux with one command (`mkdir -p {dir1,dir2,...}`): <https://www.howtogeek.com/275069/how-to-create-multiple-subdirectories-with-one-linux-command/>
- Only listing subdirectories in a directory in Linux with `ls -d */:` <https://www.cyberciti.biz/faq/linux-list-just-directories-or-directory-names/>
- Number rows in a group in increasing order in pandas: <https://stackoverflow.com/questions/37997668/pandas-number-rows-within-group-in-increasing-order>
- List only files in the current directory with `os.listdir()` and `os.path.isfile()`: <https://stackoverflow.com/question/s/11968976/list-files-only-in-the-current-directory>
- Download files from Google Colab to Google Drive with `from google.colab import files:` <https://predictivehacks.com/?all-tips=how-to-download-files-and-folders-from-colab>

Handling Imbalanced Data

- Calculating class weights to use when modeling with imbalanced data: https://www.tensorflow.org/tutorials/structured_data/imbalanced_data#train_a_model_with_class_weights
- Creating power labels and function to over/undersample classes to balance the data: <https://medium.com/the-owl/imbalanced-multilabel-image-classification-using-keras-fbd8c60d7a4b>

Multi-label classification websites

- Classification on an imbalanced dataset: https://www.tensorflow.org/tutorials/structured_data/imbalanced_data
- Multi-label classification with NN keras: <https://towardsdatascience.com/multi-label-image-classification-with-neural-network-keras-ddc1ab1afede>
- Multi-label classification in tensorflow: <https://towardsdatascience.com/multi-label-image-classification-in-tensorflow-2-0-7d4cf8a4bc72>
- **Multi-label classification in tensorflow: <https://medium.com/deep-learning-with-keras/how-to-solve-multi-label-classification-problems-in-deep-learning-with-tensorflow-keras-7fb933243595>
- Video version of above article: <https://www.youtube.com/watch?v=5MQ63pDxULw>
- Google Colab notebook: https://colab.research.google.com/drive/1dpjRslAA15q4tswH_9j8WlQ_Ro0BITSp?usp=sharing#scrollTo=Q3iCuaVXNLhS

- tf.data.Dataset API: https://www.tensorflow.org/api_docs/python/tf/data/Dataset
 - `from_tensor_slices`
- tf.io.decode_jpeg: https://www.tensorflow.org/api_docs/python/tf/io/decode_jpeg
- uint8: https://docs.microsoft.com/en-us/openspecs/windows_protocols/ms-dtyp/a88ed362-a905-4ed2-85f5-cfc8692c9842#:~:text=A%20UINT8%20is%20an%20,is%20not%20reserved%20for%20signing
- tf.io.read_file: https://www.tensorflow.org/api_docs/python/tf/io/read_file
- tf.py_function: https://www.tensorflow.org/api_docs/python/tf/py_function
- How to use the tf.data.Dataset.map() function in TensorFlow: <https://www.qcptutorials.com/article/how-to-use-map-function-with-tensorflow-datasets>
- Prefetching: https://www.tensorflow.org/guide/data_performance#prefetching
- numpy.squeeze: <https://numpy.org/doc/stable/reference/generated/numpy.squeeze.html>
- Combinations of activation and loss functions in different situations: <https://towardsdatascience.com/deep-learning-which-loss-and-activation-functions-should-i-use-ac02f1c56aa8>

Visualization

- Convert scalar tensor to a scalar variable with `.numpy()`: <https://stackoverflow.com/questions/37049411/tensorflow-how-to-convert-scalar-tensor-to-scalar-variable-in-python>
- Outline for assignment 2: M3.4 Assignment 2: Designing CNN Architecture to recognize UVA Historical Landmarks.ipynb

Creating a Convolutional Neural Network (CNN)

- Géron, Aurélien, Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems (O'Reilly Media, 2019) (for plotting images)
 - Tutorial 2.6 - Introduction to Artificial Neural Networks with Keras

Ensemble Models

- Workflow for normalizing and testing weights for ensemble models: <https://machinelearningmastery.com/weighted-average-ensemble-for-deep-learning-neural->

[networks/#:~:text=A%20weighted%20ensemble%20is%20an,the%20performanc%20of%20the%20model](#)

- `itertools.product(*iterables, repeat=1)`: <https://docs.python.org/3/library/itertools.html#itertools.product>
- Element-wise multiplication of two lists with `zip`: <https://stackoverflow.com/questions/10271484/how-to-perform-element-wise-multiplication-of-two-lists>
- `tf.keras.layers.Rescaling` to rescale pixel values within neural network: https://www.tensorflow.org/api_docs/python/tf/keras/layers/Rescaling
- Creating an ensemble model from pre-trained models: <https://www.tensorflow.org/guide/keras/functional>
- Renaming a model with `model._name = "NAME"` to prevent errors about models being named the same thing in ensemble model: <https://stackoverflow.com/questions/56886442/error-when-trying-to-rename-a-pretrained-model-on-tf-keras>

GitHub

- Creating a personal access token (instead of using a password): <https://docs.github.com/en/authentication/keeping-your-account-and-data-secure/creating-a-personal-access-token>
- Git Large File Storage: <https://git-lfs.github.com/>
 - **Steps**
 - Dowload v3.1.4 (mac - Intel Silicon)
 - `brew install git-lfs`
 - `git lfs install`
 - `cd ../Desktop/repos`
 - `git clone https://github.com/cew4pf/DS6050-eye-project.git`
 - `cd DS6050-eye-project/`
 - `git lfs track "*.zip"`
 - `git add .gitattributes`
 - `git add test_dir.zip`
 - `git commit -m "Zipped data created with Google Colab"`
 - `git push origin main`